

# [COM6513] Assignment 1: Sentiment Analysis with Logistic Regression

## Instructor: Nikos Aletras

The goal of this assignment is to develop and test a **text classification** system for **sentiment analysis**, in particular to predict the sentiment of movie reviews, i.e. positive or negative (binary classification).

For that purpose, you will implement:

- Text processing methods for extracting Bag-Of-Word features, using
  - n-grams (BOW), i.e. unigrams, bigrams and trigrams to obtain vector representations of documents where  $n=1,2,3$  respectively. Two vector weighting schemes should be tested: (1) raw frequencies (**1 mark**); (2) tf.idf (**1 mark**).
  - character n-grams (BOCN). A character n-gram is a contiguous sequence of characters given a word, e.g. for  $n=2$ , 'coffee' is split into {'co', 'of', 'ff', 'fe', 'ee'}. Two vector weighting schemes should be tested: (1) raw frequencies (**1 mark**); (2) tf.idf (**1 mark**). **Tip: Note the large vocabulary size!**
  - a combination of the two vector spaces (n-grams and character n-grams) choosing your best performing weighting respectively (i.e. raw or tfidf). (**1 mark**) **Tip: you should merge the two representations**
- Binary Logistic Regression (LR) classifiers that will be able to accurately classify movie reviews trained with:
  - (1) BOW-count (raw frequencies)
  - (2) BOW-tfidf (tf.idf weighted)
  - (3) BOCN-count
  - (4) BOCN-tfidf
  - (5) BOW+BOCN (best performing weighting; raw or tfidf)
- The Stochastic Gradient Descent (SGD) algorithm to estimate the parameters of your Logistic Regression models. Your SGD algorithm should:
  - Minimise the Binary Cross-entropy loss function (**1 mark**)
  - Use L2 regularisation (**1 mark**)
  - Perform multiple passes (epochs) over the training data (**1 mark**)
  - Randomise the order of training data after each pass (**1 mark**)
  - Stop training if the difference between the current and previous development loss is smaller than a threshold (**1 mark**)
  - After each epoch print the training and development loss (**1 mark**)
- Discuss how did you choose hyperparameters (e.g. learning rate and regularisation strength) for each LR model? You should use a table showing model performance using different set of hyperparameter values. (**2 marks**). **\*\*Tip: Instead of using all possible combinations, you could perform a random sampling of combinations.**
- After training each LR model, plot the learning process (i.e. training and validation loss in each epoch) using a line plot. Does your model underfit, overfit or is it about right? Explain why. (**1 mark**).
- Identify and show the most important features (model interpretability) for each class (i.e. top-10 most positive and top-10 negative weights). Give the top 10 for each class and comment on whether they make sense (if they don't you might have a bug!). If you were to apply the classifier into a different domain such

laptop reviews or restaurant reviews, do you think these features would generalise well? Can you propose what features the classifier could pick up as important in the new domain? (**2 marks**)

- Provide well documented and commented code describing all of your choices. In general, you are free to make decisions about text processing (e.g. punctuation, numbers, vocabulary size) and hyperparameter values. We expect to see justifications and discussion for all of your choices (**2 marks**).
- Provide efficient solutions by using Numpy arrays when possible (you can find tips in Lab 1 sheet). Executing the whole notebook with your code should not take more than 5 minutes on a any standard computer (e.g. Intel Core i5 CPU, 8 or 16GB RAM) excluding hyperparameter tuning runs (**2 marks**).

## Data

The data you will use are taken from here: <http://www.cs.cornell.edu/people/pabo/movie-review-data/> (<http://www.cs.cornell.edu/people/pabo/movie-review-data/>) and you can find it in the `./data_sentiment` folder in CSV format:

- `data_sentiment/train.csv` : contains 1,400 reviews, 700 positive (label: 1) and 700 negative (label: 0) to be used for training.
- `data_sentiment/dev.csv` : contains 200 reviews, 100 positive and 100 negative to be used for hyperparameter selection and monitoring the training process.
- `data_sentiment/test.csv` : contains 400 reviews, 200 positive and 200 negative to be used for testing.

## Submission Instructions

You should submit a Jupyter Notebook file (assignment1.ipynb) and an exported PDF version (you can do it from Jupyter: File->Download as->PDF via Latex or you can print it as PDF using your browser).

You are advised to follow the code structure given in this notebook by completing all given funtions. You can also write any auxilliary/helper functions (and arguments for the functions) that you might need but note that you can provide a full solution without any such functions. Similarly, you can just use only the packages imported below but you are free to use any functionality from the [Python Standard Library](https://docs.python.org/2/library/index.html) (<https://docs.python.org/2/library/index.html>), NumPy, SciPy (excluding built-in softmax funtcions) and Pandas. You are not allowed to use any third-party library such as Scikit-learn (apart from metric functions already provided), NLTK, Spacy, Keras etc..

There is no single correct answer on what your accuracy should be, but correct implementations usually achieve F1-scores around 80% or higher. The quality of the analysis of the results is as important as the accuracy itself.

This assignment will be marked out of 20. It is worth 20% of your final grade in the module.

The deadline for this assignment is **23:59 on Mon, 14 Mar 2022** and it needs to be submitted via Blackboard. Standard departmental penalties for lateness will be applied. We use a range of strategies to **detect unfair means** (<https://www.sheffield.ac.uk/ssid/unfair-means/index>), including Turnitin which helps detect plagiarism. Use of unfair means would result in getting a failing grade.

In [1]:

```

import pandas as pd
import numpy as np
from collections import Counter
import re
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import random

# fixing random seed for reproducibility
random.seed(123)
np.random.seed(123)

```

## Load Raw texts and labels into arrays

First, you need to load the training, development and test sets from their corresponding CSV files (tip: you can use Pandas dataframes).

In [2]:

```

dev_data = pd.read_csv('data_sentiment/dev.csv')      # Reading CSV file
dev_data.columns=['Text', 'Label']                   # Column names are given here
test_data = pd.read_csv('data_sentiment/test.csv')
test_data.columns = ['Text', 'Label']
train_data = pd.read_csv('data_sentiment/train.csv')
train_data.columns = ['Text', 'Label']

```

If you use Pandas you can see a sample of the data.

In [3]:

```
dev_data.head()
```

Out[3]:

	Text	Label
0	wong kar-wei's " fallen angels " is , on a pur...	1
1	there is nothing like american history x in th...	1
2	an unhappy italian housewife , a lonely waiter...	1
3	when people are talking about good old times ,...	1
4	the rocky horror picture show 'special edition...	1

The next step is to put the raw texts into Python lists and their corresponding labels into NumPy arrays:

In [4]:

```
dev_text = list(dev_data['Text'])           # Seprating text and Label data
dev_label = list(dev_data['Label'])
train_text = list(train_data['Text'])
train_label = list(train_data['Label'])
test_text = list(test_data['Text'])
test_label = list(test_data['Label'])
```

## Vector Representations of Text

To train and test Logisitic Regression models, you first need to obtain vector representations for all documents given a vocabulary of features (unigrams, bigrams, trigrams).

## Text Pre-Processing Pipeline

To obtain a vocabulary of features, you should:

- tokenise all texts into a list of unigrams (tip: using a regular expression)
- remove stop words (using the one provided or one of your preference)
- compute bigrams, trigrams given the remaining unigrams (or character ngrams from the unigrams)
- remove ngrams appearing in less than K documents
- use the remaining to create a vocabulary of unigrams, bigrams and trigrams (or character n-grams). You can keep top N if you encounter memory issues.

In [5]:

```
stop_words = ['a', 'in', 'on', 'at', 'and', 'or',
              'to', 'the', 'of', 'an', 'by',
              'as', 'is', 'was', 'were', 'been', 'be',
              'are', 'for', 'this', 'that', 'these', 'those', 'you', 'i',
              'it', 'he', 'she', 'we', 'they', 'will', 'have', 'has',
              'do', 'did', 'can', 'could', 'who', 'which', 'what',
              'his', 'her', 'they', 'them', 'from', 'with', 'its']
```

## N-gram extraction from a document

You first need to implement the `extract_ngrams` function. It takes as input:

- `x_raw` : a string corresponding to the raw text of a document
- `ngram_range` : a tuple of two integers denoting the type of ngrams you want to extract, e.g. (1,2) denotes extracting unigrams and bigrams.
- `token_pattern` : a string to be used within a regular expression to extract all tokens. Note that data is already tokenised so you could opt for a simple white space tokenisation.
- `stop_words` : a list of stop words
- `vocab` : a given vocabulary. It should be used to extract specific features.
- `char_ngrams` : boolean. If true the function extracts character n-grams

and returns:

- `x`: a list of all extracted features.

See the examples below to see how this function should work.

In [6]:

```

def extract_ngrams(x_raw, ngram_range=(1,3), token_pattern=r'\b[A-Za-z]{2,}\b',
                   stop_words= stop_words, vocab=None, char_ngrams = True):
    if char_ngrams == False:
        tokens = []
        for word in re.findall(token_pattern,x_raw):
            if word.lower() not in stop_words:
                tokens.append(word.lower())

        ngrams_list = []

        # Extracting tokens by words
        for num in range(0, len(tokens)):
            ngram = ' '.join(tokens[num:num + 1])
            ngrams_list.append(ngram)

        for num in range(0, len(tokens)):
            ngram = ' '.join(tokens[num:num + 2])
            ngrams_list.append(ngram)

        for num in range(0, len(tokens)):
            ngram = ' '.join(tokens[num:num + 3])
            ngrams_list.append(ngram)

        #return ngrams_list

        x = set(ngrams_list)

        # Extracting tokens by characters
        return list(x)
    elif char_ngrams == True:
        ngrams_list = []
        for x in range(len(x_raw)):
            n=x_raw[x:x+2]
            ngrams_list.append(n)
        for x in range(len(x_raw)):
            n=x_raw[x:x+3]
            ngrams_list.append(n)
        for x in range(len(x_raw)):
            n=x_raw[x:x+4]
            ngrams_list.append(n)
        x = set(ngrams_list)

        return list(x)

```

Note that it is OK to represent n-grams using lists instead of tuples: e.g. ['great', ['great', 'movie']]

For extracting character n-grams the function should work as follows:

In [7]:

```
extract_ngrams("movie",
               ngram_range=(2,4),
               stop_words=[],
               char_ngrams=True)
```

Out[7]:

```
['mov', 'movi', 'vie', 'ie', 'ovie', 'vi', 'e', 'ov', 'mo', 'ovi']
```

## Create a vocabulary

The `get_vocab` function will be used to (1) create a vocabulary of ngrams; (2) count the document frequencies of ngrams; (3) their raw frequency. It takes as input:

- `X_raw` : a list of strings each corresponding to the raw text of a document
- `ngram_range` : a tuple of two integers denoting the type of ngrams you want to extract, e.g. (1,2) denotes extracting unigrams and bigrams.
- `token_pattern` : a string to be used within a regular expression to extract all tokens. Note that data is already tokenised so you could opt for a simple white space tokenisation.
- `stop_words` : a list of stop words
- `min_df` : keep ngrams with a minimum document frequency.
- `keep_topN` : keep top-N more frequent ngrams.

and returns:

- `vocab` : a set of the n-grams that will be used as features.
- `df` : a Counter (or dict) that contains ngrams as keys and their corresponding document frequency as values.
- `ngram_counts` : counts of each ngram in vocab

Hint: it should make use of the `extract_ngrams` function.

In [8]:

```
def get_vocab(X_raw,
              ngram_range=(1, 3),
              token_pattern=r'\b[A-Za-z]{2,}\b',
              min_df=1,
              keep_topN=0,
              stop_words=stop_words):

    doc_freq = Counter()
    ngram_count = Counter()
    z=[]
    for text in X_raw:
        # A list of ngrams for the given document `text`
        ngram_list = extract_ngrams(text, ngram_range, token_pattern, stop_words, char_ngram
        doc_freq.update(set(ngram_list)) # Here we are counting docu
        for ngram in ngram_list: # Here we are counting ngra
            list=[]
            if doc_freq[ngram]>=min_df:
                list.append(ngram)
            ngram_count.update(list)
    vocab = {ngram for ngram, _ in ngram_count.most_common(keep_topN)} # Here we are extr
    return vocab, doc_freq, ngram_count
```

Now you should use `get_vocab` to create your vocabulary and get document and raw frequencies of n-grams:

In [9]:

```
vocab, doc_freq = get_vocab(train_text, keep_topN=5000)[0:2]
```

In [10]:

```
len(vocab)
```

Out[10]:

```
5000
```

In [11]:

vocab

Out[11]:

```
{'troopers',
 'police officer',
 'while',
 'san',
 'well but',
 'talented',
 'dead',
 'george clooney',
 'bell',
 'all film',
 'benefit',
 'love story',
 'upon',
 'beating',
 'never',
 'intricate',
 'inevitable',
 'wonders'.
```

Then, you need to create 2 dictionaries: (1) vocabulary id -> word; and (2) word -> vocabulary id so you can use them for reference:

In [12]:

```
id2word = {index: value for index, value in enumerate(vocab)}      # vocabulary id -> word
word_to_vocab_id = {k : v for k, v in id2word.items()}           # word -> vocabulary id
print(word_to_vocab_id)
```

```
{0: 'troopers', 1: 'police officer', 2: 'while', 3: 'san', 4: 'well but',
 5: 'talented', 6: 'dead', 7: 'george clooney', 8: 'bell', 9: 'all film', 1
0: 'benefit', 11: 'love story', 12: 'upon', 13: 'beating', 14: 'never', 1
5: 'intricate', 16: 'inevitable', 17: 'wonders', 18: 'scenario', 19: 'proj
ect', 20: 'taken', 21: 'exact', 22: 'death', 23: 'unpleasant', 24: 'anna',
25: 'there also', 26: 'finds', 27: 'clean', 28: 'my favorite', 29: 'jennif
er', 30: 'lone', 31: 'painting', 32: 'murders', 33: 'provoking', 34: 'sign
ificance', 35: 'walter', 36: 'cop', 37: 'rich', 38: 'moments', 39: 'suspen
se', 40: 'first rate', 41: 'best known', 42: 'makes up', 43: 'more', 44:
'big', 45: 'must', 46: 'marry', 47: 'profound', 48: 'timing', 49: 'positiv
e', 50: 'so many', 51: 'digital', 52: 'neve campbell', 53: 'explored', 54:
'happily', 55: 'exception', 56: 'generally', 57: 'phenomenon', 58: 'look l
ike', 59: 'jazz', 60: 'capture', 61: 'point', 62: 'coming', 63: 'regular',
64: 'disaster', 65: 'unique', 66: 'turkey', 67: 'wonder if', 68: 'qualit
y', 69: 'forms', 70: 'covers', 71: 'seen before', 72: 'bad but', 73: 'firs
t all', 74: 'camp', 75: 'purpose', 76: 'keanu', 77: 'audiences', 78: 'kean
u reeves', 79: 'something', 80: 'among', 81: 'needed', 82: 'luck', 83: 'he
lps', 84: 'training', 85: 'missed', 86: 'daniel', 87: 'supposedly', 88: 'i
nternational', 89: 'when first', 90: 'five years', 91: 'plot', 92: 'indepe
ndent', 93: 'lost', 94: 'make sure', 95: 'happily', 96: 'international', 9
```

Now you should be able to extract n-grams for each text in the training, development and test sets:



In [13]:

```
def generator_train ():
    return (extract_ngrams(text, vocab=vocab,char_ngrams=False)for text in train_text)
```

In [14]:

```
def generator_dev ():
    return (extract_ngrams(text, vocab=vocab,char_ngrams=False)for text in dev_text)
```

In [15]:

```
def generator_test ():
    return (extract_ngrams(text, vocab=vocab,char_ngrams=False)for text in test_text)
```

In [16]:

```
train_texts_ngrams = generator_train()
dev_texts_ngrams = generator_dev ()
test_texts_ngrams = generator_test()
```

## Vectorise documents

Next, write a function `vectoriser` to obtain Bag-of-ngram representations for a list of documents. The function should take as input:

- `X_ngram` : a list of texts (documents), where each text is represented as list of n-grams in the `vocab`
- `vocab` : a set of n-grams to be used for representing the documents

and return:

- `X_vec` : an array with dimensionality  $N \times |\text{vocab}|$  where  $N$  is the number of documents and  $|\text{vocab}|$  is the size of the vocabulary. Each element of the array should represent the frequency of a given n-gram in a document.

In [17]:

```
def vectorise(X_ngram, vocab):
    X_vec = []
    for ngram_list in X_ngram:
        counter = Counter(ngram_list)
        list = []
        for v in vocab:
            list.append(counter[v])
        X_vec.append(list)
    return np.array(X_vec)
```

Finally, use `vectorise` to obtain document vectors for each document in the train, development and test set. You should extract both count and tf.idf vectors respectively:

### Count vectors

In [18]:

```
train_count = vectorise(train_texts_ngrams, vocab)
dev_count = vectorise(dev_texts_ngrams, vocab)
test_count = vectorise(test_texts_ngrams, vocab)
```

In [22]:

```
Train_norm = np.log10(1 + train_count) # Reference ---> Lecture Notes
Dev_norm = np.log10(1 + dev_count)    # squash the raw frequency, by using the log10.
Test_norm = np.log10(1 + test_count)
```

In [23]:

```
train_tfidf = Train_norm * train_idf # Calculating Tfidf
dev_tfidf = Dev_norm * dev_idf      # tfidf = tf * idf
test_tfidf = Test_norm * test_idf
```

## Binary Logistic Regression

After obtaining vector representations of the data, now you are ready to implement Binary Logistic Regression for classifying sentiment.

First, you need to implement the `sigmoid` function. It takes as input:

- `z` : a real number or an array of real numbers

and returns:

- `sig` : the sigmoid of `z`

In [24]:

```
def sigmoid(z): # Reference
    sig = 1 / (1 + np.exp(-z)) # https://towardsdatascience.com/building-a-logistic-regr
    return sig
```

Then, implement the `predict_proba` function to obtain prediction probabilities. It takes as input:

- `X` : an array of inputs, i.e. documents represented by bag-of-ngram vectors ( $N \times |vocab|$ )
- `weights` : a 1-D array of the model's weights ( $1, |vocab|$ )

and returns:

- `preds_proba` : the prediction probabilities of `X` given the weights

In [25]:

```
def predict_proba(X, weights): # Reference Logic
    z = X.dot(weights) # https://pyimagesearch.com/2016/10/17/stochastic-gradient-
    preds_proba = sigmoid(z)

    return preds_proba
```

Then, implement the `predict_class` function to obtain the most probable class for each vector in an array of input vectors. It takes as input:

- `X` : an array of documents represented by bag-of-ngram vectors ( $N \times |vocab|$ )

- `weights` : a 1-D array of the model's weights ( $1, |vocab|$ )

and returns:

- `preds_class` : the predicted class for each  $x$  in  $X$  given the weights

In [26]:

```
def predict_class(X, weights):
    list = []
    for prob in predict_proba(X, weights):      # Assigning range if <= 0.5 then assign to
        if prob <= 0.5:
            list.append(0)
        else:
            list.append(1)
    preds_class = list
    return preds_class
```

To learn the weights from data, we need to minimise the binary cross-entropy loss. Implement `binary_loss` that takes as input:

- $X$  : input vectors
- $Y$  : labels
- `weights` : model weights
- `alpha` : regularisation strength

and return:

- $l$  : the loss score

In [27]:

```
def binary_loss(X, Y, weights, alpha=0.00001):
    l = -Y * np.log(predict_proba(X, weights)) - (1 - Y) * np.log(1 - predict_proba(X, weights))

    # L2 Regularisation      # Reference
    l = l + alpha * weights.dot(weights)      # https://github.com/akashmantry/LogisticRegression
                                              # Lreg = L + αR(w) --> Lecture notes

    # Return the average loss
    return np.mean(l)
```

Now, you can implement Stochastic Gradient Descent to learn the weights of your sentiment classifier. The `SGD` function takes as input:

- $X_{tr}$  : array of training data (vectors)
- $Y_{tr}$  : labels of  $X_{tr}$
- $X_{dev}$  : array of development (i.e. validation) data (vectors)
- $Y_{dev}$  : labels of  $X_{dev}$
- `lr` : learning rate
- `alpha` : regularisation strength
- `epochs` : number of full passes over the training data
- `tolerance` : stop training if the difference between the current and previous validation loss is smaller than a threshold

- `print_progress` : flag for printing the training progress (train/validation loss)

and returns:

- `weights` : the weights learned
- `training_loss_history` : an array with the average losses of the whole training set after each epoch
- `validation_loss_history` : an array with the average losses of the whole development set after each epoch

In [28]:

```
def SGD(X_tr, Y_tr, X_dev=[], Y_dev=[], lr=0.1, alpha=0.00001, epochs=5, tolerance=0.0001,
        np.random.seed(123)                                # Random seed is fixed here so that we can get sa
        training_loss_history = []
        validation_loss_history = []

        weights_int_list = []
        for i in range(train_count.shape[1]):
            weights_int_list.append(0)                        # Creating weights with all zeros so that we can
        weights_int = np.array(weights_int_list)
        weights = weights_int.astype(np.float)

        def zipper(X_tr, Y_tr):
            size = len(X_tr) if len(X_tr) < len(Y_tr) else len(Y_tr)
            retList = []                                       # Create training tuples
            for i in range(size):                               # Adding values from two list simultaneously
                retList.append((X_tr[i], Y_tr[i]))
            return retList

        train_docs = zipper(X_tr, Y_tr)

        for epoch in range(epochs):
            np.random.shuffle(train_docs)                      # Shuffling to randomise all values
                                                                # Reference
            for first, second in train_docs:                   #  $w = w - \eta \nabla_w L(w; x_i; y_i)$  --> Lecture Notes
                weights = weights - lr * (first * (predict_proba(first, weights) - second) + 2

            # Monitor training and validation loss
            loss_in_training = binary_loss(X_tr, Y_tr, weights, alpha)
            loss_in_dev = binary_loss(X_dev, Y_dev, weights, alpha)

            # Early stopping                                     # Reference
                                                                # previous validation loss - current validation loss; s
            if epoch > 0 and validation_loss_history[-1] - loss_in_dev < tolerance:
                break
            else:
                training_loss_history.append(loss_in_training)
                validation_loss_history.append(loss_in_dev)

            if print_progress:
                print("Epoch:- ", epoch, " ", "Training loss:- ", loss_in_training, " ", "Validatio

        return weights, training_loss_history, validation_loss_history
```

## Train and Evaluate Logistic Regression with Count vectors

First train the model using SGD:

In [29]:

```
w_count, training_loss_count, dev_loss_count = SGD(X_tr=train_count, Y_tr=np.array(train_lab
```

```
<ipython-input-28-1bbe47e1eeaf>:11: DeprecationWarning: `np.float` is a de  
precated alias for the builtin `float`. To silence this warning, use `floa  
t` by itself. Doing this will not modify any behavior and is safe. If you  
specifically wanted the numpy scalar type, use `np.float64` here.  
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org  
/devdocs/release/1.20.0-notes.html#deprecations (https://numpy.org/devdoc  
s/release/1.20.0-notes.html#deprecations)  
    weights = weights_int.astype(np.float)
```

Now plot the training and validation history per epoch for the best hyperparameter combination. Does your model underfit, overfit or is it about right? Explain why.

In [30]:

```
plt.plot(training_loss_count, label='Train_loss')
plt.plot(dev_loss_count, label='Valid_loss')

plt.xlabel('Epochs',fontweight='bold')
plt.ylabel('Loss',fontweight='bold')

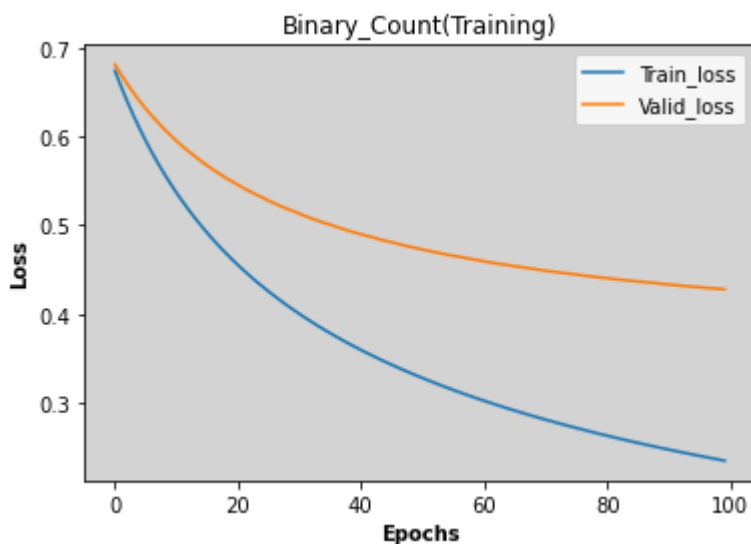
plt.title('Binary_Count(Training)')
ax = plt.axes()
ax.set_facecolor("lightgray")

plt.legend()

plt.show()
```

<ipython-input-30-41ab2a87d266>:8: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

```
ax = plt.axes()
```



Explain here...

From the plot above, it is clear that.. (i) Train loss is decreasing as count of epoch increases until it reaches a certain standard of stability..(ii) Valid loss is also decreasing as count of epoch increases until it reaches a certain standard of stability. From this I conclude that, the model is about right.

## Evaluation

Compute accuracy, precision, recall and F1-scores:

In [31]:

```
array_test_label = np.array(test_label) # Changing list to array as accu
```

In [32]:

```
preds_te_count = predict_class(test_count, w_count)

print('Accuracy:', accuracy_score(array_test_label,preds_te_count))
print('Precision:', precision_score(array_test_label,preds_te_count))
print('Recall:', recall_score(array_test_label,preds_te_count))
print('F1-Score:', f1_score(array_test_label,preds_te_count))
```

```
Accuracy: 0.8596491228070176
Precision: 0.8454106280193237
Recall: 0.8793969849246231
F1-Score: 0.8620689655172413
```

Finally, print the top-10 words for the negative and positive class respectively.

In [33]:

```
top_neg = w_count.argsort()[::-10]      # Printing top ten negative elements
for i in top_neg:
    print(id2word[i])
```

```
bad
script
worst
unfortunately
nothing
plot
boring
only
looks
supposed
```

In [34]:

```
top_pos = w_count.argsort()[::-1][:10]  # Printing top ten positive elements
for i in top_pos:
    print(id2word[i])
```

```
hilarious
also
both
great
well
many
seen
true
perfect
perfectly
```

If we were to apply the classifier we've learned into a different domain such laptop reviews or restaurant reviews, do you think these features would generalise well? Can you propose what features the classifier could pick up as important in the new domain?

The classifier here predicts a good set of positive and negative words and it is expected that it will correctly predict the some of the reviews about laptops and restaurant. However,specially for laptop and restaurant there may have some words which defines best positive and best negative words for review and using this classifier



will give accuracy but not to a great extent and it will somewhat make model underfit. Hence I don't think this features would generalise well in laptop and restaurant reviews.

## Discuss how did you choose model hyperparameters (e.g. learning rate and regularisation strength)? What is the relation between training epochs and learning rate? How the regularisation strength affects performance?

I have carried out trial and error strategy to choose model hyperparameter. As mentioned in the instruction, *Instead of using all possible combinations, you could perform a random sampling of combinations.* I have done the same. In hyperparameter optimisation, the lower bound and upper bound values are need to be defined where lower bound value is the baseline for optimal performance of model and upper bound is the one where the performance starts to mitigate. For this purpose I have set the initial lower and upper bound for learning rate as 0.0001 and 0.1 where for regularisation I have set it to 0.00001 and 0.01 precisely. After that I started changing the lower bound and upper bound to see the results in terms of precision, recall and F1-score.

### Count Vectors

Table showing model performance for learning rate :-

Trial	Learning rate	Epochs	Tr. loss	Val. loss	Precision	Recall	F1-Score
0	0.0001	99	0.2033	0.4097	0.8523	0.876	0.8634
1	0.00011	99	0.1945	0.4086	0.8522	0.868	0.8567
2	0.000105	99	0.1975	0.4083	0.8457	0.871	0.8592
3	0.00010124	99	0.2033	0.4097	0.8523	0.876	0.8635

Table showing model performance for Regularisation strength :-

Trial	Alpha	Epochs	Tr. loss	Val. loss	Precision	Recall	F1-Score
0	0.00001	99	0.2031	0.40954	0.8523	0.874	0.8640
1	0.00002	99	0.2017	0.40953	0.8522	0.873	0.8640
2	0.000015	99	0.2022	0.40948	0.8522	0.874	0.8639
3	0.0000124	99	0.2029	0.40944	0.8523	0.874	0.8640

### TF.IDF Vectors

Table showing model performance for learning rate :-

Trial	Learning rate	Epochs	Tr. loss	Val. loss	Precision	Recall	F1-Score
0	0.0001	23	0.5025	0.5845	0.8412	0.867	0.8654
1	0.0002	23	0.4053	0.5321	0.8632	0.875	0.8733
2	0.0003	23	0.3439	0.4998	0.8654	0.848	0.8798
3	0.0025	23	0.0856	0.3750	0.8869	0.875	0.8759

Table showing model performance for Regularisation strength :-

Trial	Alpha	Epochs	Tr. loss	Val. loss	Precision	Recall	F1-Score
0	0.00001	23	0.0891	0.3756	0.8878	0.873	0.8810

Trial	Alpha	Epochs	Tr. loss	Val. loss	Precision	Recall	F1-Score
1	0.00002	23	0.0903	0.3785	0.8880	0.873	0.8810
2	0.0004	23	0.1645	0.4265	0.88	0.88	0.8866
3	0.0008	23	0.2206	0.4583	0.875	0.875	0.8860

## Relationship Between Epochs and Learning Rate

The relationship between Epochs and learning rate follows a proportionality rule, where greater the learning rate, larger the weight update after each epochs.

## How regularisation strength affects performance.

Regularisation rate affects the convergence of epoch which in turn affects the model performance, so we can say that regularisation strength indirectly correlates with the model performance.

## Train and Evaluate Logistic Regression with TF.IDF vectors

Follow the same steps as above (i.e. evaluating count n-gram representations).

**Now repeat the training and evaluation process for BOW-tfidf, BOCN-count, BOCN-tfidf, BOW+BOCN including hyperparameter tuning for each model...**

In [35]:

```
w_tfidf, training_loss_tfidf, dev_loss_tfidf = SGD(X_tr=train_tfidf,Y_tr=np.array(train_lab
```

```
Epoch:- 0    Training loss:- 0.6619948876678392    Validation loss:- 0.
6708951556266846
Epoch:- 1    Training loss:- 0.6344381228181969    Validation loss:- 0.
651831140695691
Epoch:- 2    Training loss:- 0.6097707917410218    Validation loss:- 0.
6342512289210034
Epoch:- 3    Training loss:- 0.5877746679357756    Validation loss:- 0.
619360059437997
```

```
<ipython-input-28-1bbe47e1eeaf>:11: DeprecationWarning: `np.float` is a de
precated alias for the builtin `float`. To silence this warning, use `floa
t` by itself. Doing this will not modify any behavior and is safe. If you
specifically wanted the numpy scalar type, use `np.float64` here.
```

```
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations (https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations)
```

```
weights = weights_int.astype(np.float)
```

In [36]:

```
plt.plot(training_loss_tfidf, label='Train_loss')
plt.plot(dev_loss_tfidf, label='Valid_loss')

plt.xlabel('Epochs',fontweight='bold')
plt.ylabel('Loss',fontweight='bold')

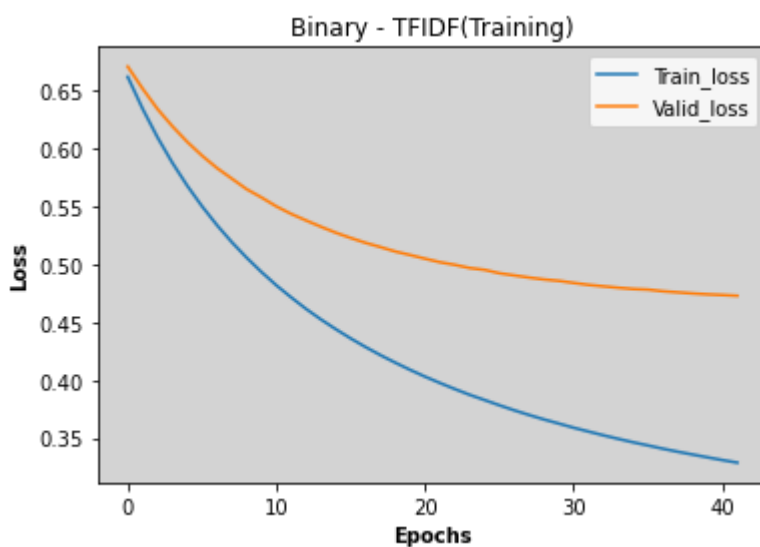
plt.title('Binary - TFIDF(Training)')
ax = plt.axes()
ax.set_facecolor("lightgray")

plt.legend()

plt.show()
```

<ipython-input-36-77ed45a3342a>:8: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

```
ax = plt.axes()
```



In [37]:

```
preds_te_tfidf = predict_class(test_tfidf, w_tfidf)

print('Accuracy:', accuracy_score(np.array(test_label),preds_te_tfidf))
print('Precision:', precision_score(np.array(test_label),preds_te_tfidf))
print('Recall:', recall_score(np.array(test_label),preds_te_tfidf))
print('F1-Score:', f1_score(np.array(test_label),preds_te_tfidf))
```

Accuracy: 0.8621553884711779  
Precision: 0.86  
Recall: 0.864321608040201  
F1-Score: 0.8621553884711779

In [38]:

```
top_neg = w_tfidf.argsort()[::-10]
for i in top_neg:
    print(id2word[i])
```

bad  
script  
worst  
nothing  
plot  
unfortunately  
boring  
looks  
only  
least

In [39]:

```
top_pos = w_tfidf.argsort()[::-1][:10]
for i in top_pos:
    print(id2word[i])
```

hilarious  
also  
both  
great  
many  
well  
perfect  
true  
best  
seen

## BOCN

In [40]:

```

def extract_ngrams_c(x_raw, ngram_range=(1,3), token_pattern=r'\b[A-Za-z]{2,}\b',
                      stop_words= stop_words, vocab=None, char_ngrams = True):
    if char_ngrams == False:
        tokens = []
        for word in re.findall(token_pattern,x_raw):
            if word.lower() not in stop_words:
                tokens.append(word.lower())

        ngrams_list = []

        for num in range(0, len(tokens)):
            ngram = ' '.join(tokens[num:num + 1])           # Extracting tokens by words
            ngrams_list.append(ngram)

        for num in range(0, len(tokens)):
            ngram = ' '.join(tokens[num:num + 2])
            ngrams_list.append(ngram)

        for num in range(0, len(tokens)):
            ngram = ' '.join(tokens[num:num + 3])
            ngrams_list.append(ngram)

        #return ngrams_list

        x = set(ngrams_list)

        return list(x)
    elif char_ngrams == True:
        ngrams_list = []
        for x in range(len(x_raw)):           # Extracting tokens by charact
            n=x_raw[x:x+2]
            ngrams_list.append(n)
        for x in range(len(x_raw)):
            n=x_raw[x:x+3]
            ngrams_list.append(n)
        for x in range(len(x_raw)):
            n=x_raw[x:x+4]
            ngrams_list.append(n)
        x = set(ngrams_list)

        return list(x)

```

In [41]:

```
def get_vocab_c(X_raw,
                ngram_range=(1, 3),
                token_pattern=r'\b[A-Za-z]{2,}\b',
                min_df=1,
                keep_topN=0,
                stop_words=stop_words):

    doc_freq_c = Counter()
    ngram_count_c = Counter()
    for text in X_raw:
        # A list of ngrams for the given document `text`
        ngram_list = extract_ngrams_c(text, ngram_range, token_pattern, stop_words, char_ngr
        doc_freq_c.update(set(ngram_list))      # Here we are counting document frequency
        for ngram in ngram_list:                # Here we are counting ngram frequency
            list=[]
            if doc_freq_c[ngram]>=min_df:
                list.append(ngram)
            ngram_count_c.update(list)
    vocab_c = {ngram for ngram, _ in ngram_count_c.most_common(keep_topN)} # Here we are
    return vocab_c, doc_freq_c, ngram_count_c
```

In [42]:

```
vocab_c, doc_freq_c = get_vocab_c(train_text, keep_topN=5000)[0:2]
```

In [43]:

```
len(vocab_c)
```

Out[43]:

5000

In [44]:

```
vocab_c
```

Out[44]:

```
{'t f',
 'nit',
 'lp',
 'san',
 ' cr',
 'ood ',
 'e an',
 'his',
 'y ) ',
 'ng s',
 "'s c",
 'at b',
 'ille',
 'm i',
 'o ',
 ' lon',
 'his ',
 ' who'.
```



In [50]:

```
def vectorise_c(X_ngram, vocab):
    X_vec_c = []
    for ngram_list in X_ngram:
        counter = Counter(ngram_list)
        list_c = []
        for v in vocab:
            list_c.append(counter[v])
        X_vec_c.append(list_c)
    return np.array(X_vec_c)
```

*# Here we are creating a function to vectorise t*  
*# Here we are using temporary list variable to a*

In [51]:

```
train_count_c = vectorise_c(train_texts_ngrams_c, vocab_c)

dev_count_c = vectorise_c(dev_texts_ngrams_c, vocab_c)

test_count_c = vectorise_c(test_texts_ngrams_c, vocab_c)
```

In [52]:

```
train_count_c.shape
```

Out[52]:

(1399, 5000)

In [53]:

```
total_train_docs_c = len(train_text)
total_dev_docs_c = len(dev_text)
total_test_docs_c = len(test_text)

dev_df_c = get_vocab_c(dev_text, keep_topN=5000)[1]
test_df_c = get_vocab_c(test_text, keep_topN=5000)[1]
for v in vocab_c:
    train_idf_c = np.array([np.log10(total_train_docs_c/doc_freq_c[v])])
for v in vocab_c:
    if dev_df_c[v]:
        dev_idf_c = np.array([np.log10(total_dev_docs_c / dev_df_c[v])])
for v in vocab_c:
    if test_df_c[v]:
        test_idf_c = np.array([np.log10(total_test_docs_c/test_df_c[v])])
```

*# Reference --> Lecture Notes*

In [54]:

```
Train_norm_c = np.log10(1 + train_count_c)
Dev_norm_c = np.log10(1 + dev_count_c)
Test_norm_c = np.log10(1 + test_count_c)
```

*# Reference ---> Lecture Notes*  
*# squash the raw frequency, by using the log10.*

In [55]:

```
train_tfidf_c = Train_norm_c * train_idf_c
dev_tfidf_c = Dev_norm_c * dev_idf_c
test_tfidf_c = Test_norm_c * test_idf_c
```

*# Calculating Tfidf*  
*# tfidf = tf \* idf*



In [56]:

```
def sigmoid_c(z):          # Reference
    return 1 / (1 + np.exp(-z)) # https://towardsdatascience.com/building-a-logistic-re
```

In [57]:

```
def predict_proba_c(X, weights): # Reference Logic
    z = X.dot(weights)           # https://pyimagesearch.com/2016/10/17/stochastic-gradi
    return sigmoid_c(z)
```

In [58]:

```
def predict_class_c(X, weights):
    list = []
    for prob in predict_proba_c(X, weights):
        if prob <= 0.5:
            list.append(0)          # Assigning range if <= 0.5 then assign to 0 else
        else:
            list.append(1)
    return list
```

In [59]:

```
def binary_loss_c(X, Y, weights, alpha=0.00001):
    l = -Y * np.log(predict_proba_c(X, weights)) - (1 - Y) * np.log(1 - predict_proba_c(X,

    # L2 Regularisation          # Reference
    l += alpha * weights.dot(weights) # https://github.com/akashmantry/LogisticRegression
                                     # Lreg = L + αR(w) --> Lecture notes

    # Return the average loss
    return np.mean(l)
```

In [60]:

```

def SGD_c(X_tr, Y_tr, X_dev, Y_dev, lr=0.1, alpha=0.00001, epochs=5, tolerance=0.0001, print_progress=True):
    np.random.seed(123) # Random seed is fixed here so that we can get same results every time
    training_loss_history_c = []
    validation_loss_history_c = []

    weights_int_list_c = [] # Creating weights with all zeros so that we can start from a neutral point
    for i in range(train_count_c.shape[1]):
        weights_int_list_c.append(0)
    weights_int_c = np.array(weights_int_list_c)
    weights_c = weights_int_c.astype(np.float)

    def zipper_c(X_tr, Y_tr): # Create training tuples
        size = len(X_tr) if len(X_tr) < len(Y_tr) else len(Y_tr) # Adding values from two lists
        retList = []
        for i in range(size):
            retList.append((X_tr[i], Y_tr[i]))
        return retList

    train_docs_c = zipper_c(X_tr, Y_tr)

    for epoch in range(epochs):
        np.random.shuffle(train_docs_c) # Shuffling to randomise all values
        # Reference
        for first, second in train_docs_c: #  $w = w - \eta \nabla_w L(w; x_i; y_i)$  --> Lecture Notes
            weights_c = weights_c - lr * (first * (predict_proba_c(first, weights_c) - second))

        # Monitor training and validation loss
        loss_in_training_c = binary_loss_c(X_tr, Y_tr, weights_c, alpha)
        loss_in_dev_c = binary_loss_c(X_dev, Y_dev, weights_c, alpha)

        # Early stopping # Reference
        # previous validation loss - current validation loss; should be > tolerance
        if epoch > 0 and validation_loss_history_c[-1] - loss_in_dev_c < tolerance:
            break
        else:
            training_loss_history_c.append(loss_in_training_c)
            validation_loss_history_c.append(loss_in_dev_c)

    if print_progress:
        #print(f'Epoch: {epoch} | Training loss: {cur_loss_tr} | Validation loss: {cur_loss_dev}')
        print("Epoch:- ", epoch, " ", "Training loss:- ", loss_in_training_c, " ", "Validation loss:- ", loss_in_dev_c)
    return weights_c, training_loss_history_c, validation_loss_history_c

```

In [61]:

```
w_count_c, training_loss_count_c, dev_loss_count_c = SGD_c(X_tr=train_count_c,Y_tr=np.array
```

<ipython-input-60-962918fee0ef>:11: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here.

Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations> (<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>)

```
weights_c = weights_int_c.astype(np.float)
```

In [62]:

```
plt.plot(training_loss_count_c, label='Train_loss')
plt.plot(dev_loss_count_c, label='Valid_loss')

plt.xlabel('Epochs',fontweight='bold')
plt.ylabel('Loss',fontweight='bold')

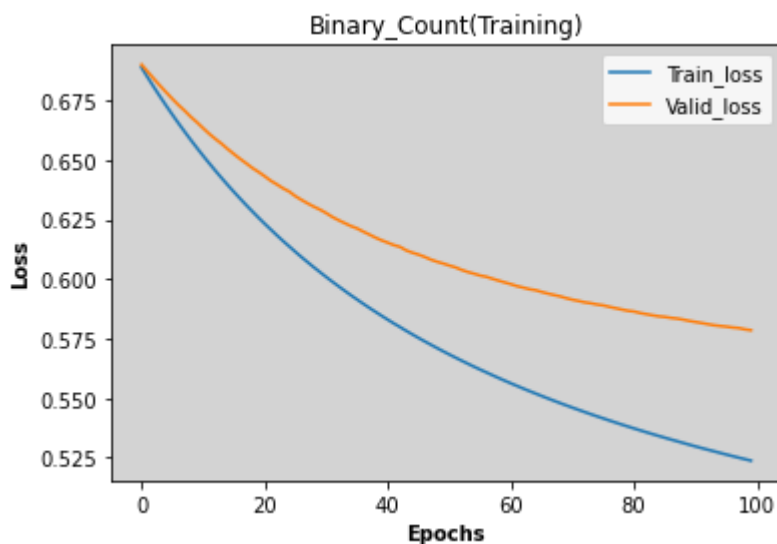
plt.title('Binary_Count(Training)')
ax = plt.axes()
ax.set_facecolor("lightgray")

plt.legend()

plt.show()
```

<ipython-input-62-4417a4b65adf>:8: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

```
ax = plt.axes()
```



In [63]:

```
array_test_label_c = np.array(test_label)    # Changing list to array as accuracy score ta
```

In [64]:

```
preds_te_count_c = predict_class_c(test_count_c, w_count_c)

print('Accuracy:', accuracy_score(array_test_label_c,preds_te_count_c))
print('Precision:', precision_score(array_test_label_c,preds_te_count_c))
print('Recall:', recall_score(array_test_label_c,preds_te_count_c))
print('F1-Score:', f1_score(array_test_label_c,preds_te_count_c))
```

```
Accuracy: 0.6992481203007519
Precision: 0.7046632124352331
Recall: 0.6834170854271356
F1-Score: 0.6938775510204082
```

In [65]:

```
w_tfidf_c, training_loss_tfidf_c, dev_loss_tfidf_c = SGD_c(X_tr=train_tfidf_c,Y_tr=np.array
```

```
Epoch:- 0    Training loss:- 0.6904966795485448    Validation loss:- 0.69
12148003649882
Epoch:- 1    Training loss:- 0.6879143906142949    Validation loss:- 0.68
93460373561778
Epoch:- 2    Training loss:- 0.685398383432307    Validation loss:- 0.687
5191950035022
Epoch:- 3    Training loss:- 0.6829470269153042    Validation loss:- 0.68
57156108230362
```

<ipython-input-60-962918fee0ef>:11: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here.  
 Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations> (<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>)

```
weights_c = weights_int_c.astype(np.float)
```

```
Epoch:- 4    Training loss:- 0.680558160644537    Validation loss:- 0.683
9878531458286
Epoch:- 5    Training loss:- 0.6782310802076481    Validation loss:- 0.68
22830053666049
Epoch:- 6    Training loss:- 0.6759636483164315    Validation loss:- 0.68
06552389948863
Epoch:- 7    Training loss:- 0.6737528936179766    Validation loss:- 0.67
91565087576937
Epoch:- 8    Training loss:- 0.6715987837294469    Validation loss:- 0.67
75782747496061
Epoch:- 9    Training loss:- 0.6695001744519493    Validation loss:- 0.67
61485316169866
Epoch:- 10   Training loss:- 0.6674506086436001    Validation loss:- 0.6
746425538482389
Epoch:- 11   Training loss:- 0.6654592925586996    Validation loss:- 0.6
731673889756857
Epoch:- 12   Training loss:- 0.6635130903135239    Validation loss:- 0.6
718266619510327
Epoch:- 13   Training loss:- 0.6616168004478585    Validation loss:- 0.6
705692210081527
Epoch:- 14   Training loss:- 0.6597692076603948    Validation loss:- 0.6
692164750684666
Epoch:- 15   Training loss:- 0.6579683188466647    Validation loss:- 0.6
679330967753333
Epoch:- 16   Training loss:- 0.6562125801024464    Validation loss:- 0.6
666896282719543
Epoch:- 17   Training loss:- 0.6544973561428005    Validation loss:- 0.6
655142793448491
Epoch:- 18   Training loss:- 0.6528302780182359    Validation loss:- 0.6
643263326142681
Epoch:- 19   Training loss:- 0.6511946014688571    Validation loss:- 0.6
632791263933235
Epoch:- 20   Training loss:- 0.64960538135053    Validation loss:- 0.662
1994711040016
Epoch:- 21   Training loss:- 0.6480596468159974    Validation loss:- 0.6
610611943364745
Epoch:- 22   Training loss:- 0.6465434471536388    Validation loss:- 0.6
600869591122408
```

Epoch:- 23	Training loss:- 0.645068576042016	Validation loss:- 0.65
91066316909521		
Epoch:- 24	Training loss:- 0.6436397865873372	Validation loss:- 0.6
582540602283568		
Epoch:- 25	Training loss:- 0.6422264631067357	Validation loss:- 0.6
571663337373143		
Epoch:- 26	Training loss:- 0.6408572043309972	Validation loss:- 0.6
562113329625406		
Epoch:- 27	Training loss:- 0.6395199364942084	Validation loss:- 0.6
553189405528105		
Epoch:- 28	Training loss:- 0.6382144288631433	Validation loss:- 0.6
544948083909319		
Epoch:- 29	Training loss:- 0.6369467834926484	Validation loss:- 0.6
537361040835109		
Epoch:- 30	Training loss:- 0.6357022008609299	Validation loss:- 0.6
528996591755573		
Epoch:- 31	Training loss:- 0.634487537703162	Validation loss:- 0.65
19781291926887		
Epoch:- 32	Training loss:- 0.6333072223571137	Validation loss:- 0.6
511708657201877		
Epoch:- 33	Training loss:- 0.6321496257458664	Validation loss:- 0.6
504381205381176		
Epoch:- 34	Training loss:- 0.631020405955202	Validation loss:- 0.64
97141497734192		
Epoch:- 35	Training loss:- 0.6299213123352827	Validation loss:- 0.6
491141699736936		
Epoch:- 36	Training loss:- 0.6288437483137984	Validation loss:- 0.6
483826365280088		
Epoch:- 37	Training loss:- 0.627792778156486	Validation loss:- 0.64
76876601753361		
Epoch:- 38	Training loss:- 0.6267697340866303	Validation loss:- 0.6
469560402377199		
Epoch:- 39	Training loss:- 0.62577109504772	Validation loss:- 0.646
2968754880194		
Epoch:- 40	Training loss:- 0.6247893995781216	Validation loss:- 0.6
457195940601848		
Epoch:- 41	Training loss:- 0.6238330308730168	Validation loss:- 0.6
451456599410135		
Epoch:- 42	Training loss:- 0.622907769239038	Validation loss:- 0.64
46865211713771		
Epoch:- 43	Training loss:- 0.621990637125021	Validation loss:- 0.64
39500472776369		
Epoch:- 44	Training loss:- 0.6211004221479085	Validation loss:- 0.6
433982353769517		
Epoch:- 45	Training loss:- 0.6202294821263465	Validation loss:- 0.6
429330341276622		
Epoch:- 46	Training loss:- 0.6193795408687143	Validation loss:- 0.6
424055070556662		
Epoch:- 47	Training loss:- 0.6185523986549033	Validation loss:- 0.6
418127534026768		
Epoch:- 48	Training loss:- 0.6177405997587774	Validation loss:- 0.6
413214054657524		
Epoch:- 49	Training loss:- 0.6169444931016976	Validation loss:- 0.6
409139392368264		
Epoch:- 50	Training loss:- 0.6161696317570117	Validation loss:- 0.6
404589995523223		
Epoch:- 51	Training loss:- 0.6154144437424934	Validation loss:- 0.6
400543657975111		
Epoch:- 52	Training loss:- 0.6146721579861727	Validation loss:- 0.6
395194951842923		
Epoch:- 53	Training loss:- 0.613950512277618	Validation loss:- 0.63

90642984891007

Epoch:- 54 Training loss:- 0.6132409817658497 Validation loss:- 0.6

386889410985886

Epoch:- 55 Training loss:- 0.6125495430507765 Validation loss:- 0.6

382842781491925

Epoch:- 56 Training loss:- 0.6118787950755831 Validation loss:- 0.6

380094188129629

Epoch:- 57 Training loss:- 0.6112159077092548 Validation loss:- 0.6

376084825497832

Epoch:- 58 Training loss:- 0.6105684468967958 Validation loss:- 0.6

372249852269722

Epoch:- 59 Training loss:- 0.6099342206403221 Validation loss:- 0.6

368385280955706

Epoch:- 60 Training loss:- 0.6093158962334234 Validation loss:- 0.6

36479997351823

Epoch:- 61 Training loss:- 0.608712076620917 Validation loss:- 0.63

60716884566442

Epoch:- 62 Training loss:- 0.6081199126310607 Validation loss:- 0.6

357828762129853

Epoch:- 63 Training loss:- 0.6075416844678508 Validation loss:- 0.6

354582779947391

Epoch:- 64 Training loss:- 0.6069815125984883 Validation loss:- 0.6

352588358636613

Epoch:- 65 Training loss:- 0.6064251797514182 Validation loss:- 0.6

34920600417333

Epoch:- 66 Training loss:- 0.6058821679186666 Validation loss:- 0.6

346028260796892

Epoch:- 67 Training loss:- 0.6053533133986634 Validation loss:- 0.6

343262217219329

Epoch:- 68 Training loss:- 0.6048355743651367 Validation loss:- 0.6

340569617137566

Epoch:- 69 Training loss:- 0.6043268551978904 Validation loss:- 0.6

337115608781594

Epoch:- 70 Training loss:- 0.6038331569017004 Validation loss:- 0.6

33406484987694

Epoch:- 71 Training loss:- 0.6033474111540452 Validation loss:- 0.6

331631692582707

Epoch:- 72 Training loss:- 0.6028733997493327 Validation loss:- 0.6

329042721094521

Epoch:- 73 Training loss:- 0.6024050587065095 Validation loss:- 0.6

327409716237511

Epoch:- 74 Training loss:- 0.6019497888711972 Validation loss:- 0.6

32525273056419

Epoch:- 75 Training loss:- 0.6015085511101819 Validation loss:- 0.6

323764032399776

Epoch:- 76 Training loss:- 0.6010696505272307 Validation loss:- 0.6

321239352822157

Epoch:- 77 Training loss:- 0.6006434009345868 Validation loss:- 0.6

319180256713917

Epoch:- 78 Training loss:- 0.600223881844529 Validation loss:- 0.63

16816557953381

Epoch:- 79 Training loss:- 0.5998148450714139 Validation loss:- 0.6

314661202546591

Epoch:- 80 Training loss:- 0.5994173299296831 Validation loss:- 0.6

313418608829497

Epoch:- 81 Training loss:- 0.599022568286133 Validation loss:- 0.63

10726372519033

Epoch:- 82 Training loss:- 0.5986391409054845 Validation loss:- 0.6

308675121232522

Epoch:- 83 Training loss:- 0.5982646586755134 Validation loss:- 0.6

306671301765195

Epoch:- 84    Training loss:- 0.5978955735054345    Validation loss:- 0.6305092030870798

In [66]:

```
plt.plot(training_loss_tfidf_c, label='Train_loss')
plt.plot(dev_loss_tfidf_c, label='Valid_loss')

plt.xlabel('Epochs',fontweight='bold')
plt.ylabel('Loss',fontweight='bold')


plt.title('Binary - TFIDF(Training)')
ax = plt.axes()
ax.set_facecolor("lightgray")

plt.legend()

plt.show()
```

<ipython-input-66-735fd7e32112>:8: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

```
ax = plt.axes()
```



In [67]:

```
preds_te_tfidf_c = predict_class_c(test_tfidf_c, w_tfidf_c)

print('Accuracy:', accuracy_score(np.array(test_label),preds_te_tfidf_c))
print('Precision:', precision_score(np.array(test_label),preds_te_tfidf_c))
print('Recall:', recall_score(np.array(test_label),preds_te_tfidf_c))
print('F1-Score:', f1_score(np.array(test_label),preds_te_tfidf_c))
```

Accuracy: 0.7042606516290727  
Precision: 0.7055837563451777  
Recall: 0.6984924623115578  
F1-Score: 0.7020202020202021



# BOW+BOCN

In [68]:

```
def extract_ngrams_c_w(x_raw, ngram_range=(1,3), token_pattern=r'\b[A-Za-z]{2,}\b', stop_wo
tokens = []
for word in re.findall(token_pattern,x_raw):
    if word.lower() not in stop_words:
        tokens.append(word.lower())

ngrams_list_c_w = []

for num in range(0, len(tokens)):
    ngram = ' '.join(tokens[num:num + 1])
    ngrams_list_c_w.append(ngram)

for num in range(0, len(tokens)):
    ngram = ' '.join(tokens[num:num + 2])
    ngrams_list_c_w.append(ngram)

for num in range(0, len(tokens)):
    ngram = ' '.join(tokens[num:num + 3])
    ngrams_list_c_w.append(ngram)

for x in range(len(x_raw)):
    n=x_raw[x:x+2]
    ngrams_list_c_w.append(n)
for x in range(len(x_raw)):
    n=x_raw[x:x+3]
    ngrams_list_c_w.append(n)
for x in range(len(x_raw)):
    n=x_raw[x:x+4]
    ngrams_list_c_w.append(n)
x = set(ngrams_list_c_w)

return list(x)
```

*# Creating tokens by words*

*# Creating tokens by characters*

*# Appending both tokens by words and characters*

In [69]:

```
def get_vocab_c_w(X_raw,
                  ngram_range=(1, 3),
                  token_pattern=r'\b[A-Za-z]{2,}\b',
                  min_df=1,
                  keep_topN=0,
                  stop_words=stop_words):

    doc_freq_c_w = Counter()           # Using counter so we can count no of times elemen
    ngram_count_c_w = Counter()        # Also counter function provides most common param
                                      # Also counter function provides update paramter t

    for text in X_raw:
        # A list of ngrams for the given document `text`
        ngram_list_c_w = extract_ngrams_c_w(text, ngram_range, token_pattern, stop_words) #
        doc_freq_c_w.update(set(ngram_list_c_w)) # Here we are using set so as to get uniq
        for ngram in ngram_list_c_w:         # Here we are counting ngram frequency
            list=[]
            if doc_freq_c_w[ngram]>=min_df:
                list.append(ngram)
            ngram_count_c_w.update(list)
    vocab_c_w = {ngram for ngram, _ in ngram_count_c_w.most_common(keep_topN)} # Here we
    return vocab_c_w, doc_freq_c_w, ngram_count_c_w
```

In [70]:

```
vocab_c_w, doc_freq_c_w = get_vocab_c_w(train_text, keep_topN=5000)[0:2]
```

In [71]:

```
len(vocab_c_w)
```

Out[71]:

5000

In [72]:

```
vocab_c_w
```

Out[72]:

```
{'t f',  
'nit',  
'lp',  
'while',  
'san',  
' cr',  
'ood ',  
'e an',  
'his',  
'y ) ',  
'ng s',  
's c",  
'at b',  
'ille',  
'm i',  
'never',  
'o ',  
' lon'.
```

In [73]:

```
def generator_train_c_w ():  
    return (extract_ngrams_c_w(text, vocab=vocab_c_w) for text in train_text)
```

In [74]:

```
def generator_dev_c_w ():  
    return (extract_ngrams_c_w(text, vocab=vocab_c_w) for text in dev_text)
```

In [75]:

```
def generator_test_c_w ():  
    return (extract_ngrams_c_w(text, vocab=vocab_c_w) for text in test_text)
```

In [76]:

```
train_texts_ngrams_c_w = generator_train()  
dev_texts_ngrams_c_w = generator_dev ()  
test_texts_ngrams_c_w = generator_test()
```



In [81]:

```

total_train_docs_c_w = len(train_text)
total_dev_docs_c_w = len(dev_text)
total_test_docs_c_w = len(test_text)

dev_df_c_w = get_vocab_c_w(dev_text, keep_topN=5000)[1]
test_df_c_w = get_vocab_c_w(test_text, keep_topN=5000)[1]
for v in vocab_c_w: # Reference --> Lecture Notes
    train_idf_c_w = np.array([np.log10(total_train_docs_c_w/doc_freq_c_w[v])])
for v in vocab_c_w:
    if dev_df_c_w[v]:
        dev_idf_c_w = np.array([np.log10(total_dev_docs_c_w / dev_df_c_w[v])])
for v in vocab_c_w:
    if test_df_c_w[v]:
        test_idf_c_w = np.array([np.log10(total_test_docs_c_w/test_df_c_w[v])])

```

In [82]:

```

Train_norm_c_w = np.log10(1 + train_count_c_w) # Reference ---> Lecture Notes
Dev_norm_c_w = np.log10(1 + dev_count_c_w) # squash the raw frequency, by using the Lo
Test_norm_c_w = np.log10(1 + test_count_c_w)

```

In [83]:

```

train_tfidf_c_w = Train_norm_c_w * train_idf_c_w # Calculating Tfidf
dev_tfidf_c_w = Dev_norm_c_w * dev_idf_c_w # tfidf = tf * idf
test_tfidf_c_w = Test_norm_c_w * test_idf_c_w

```

In [84]:

```

def sigmoid_c_w(z): # Reference
    return 1 / (1 + np.exp(-z)) # https://towardsdatascience.com/building-a-logistic-re

```

In [85]:

```

def predict_proba_c_w(X, weights): # Reference Logic
    z = X.dot(weights) # https://pyimagesearch.com/2016/10/17/stochastic-gradi
    return sigmoid_c_w(z)

```

In [86]:

```

def predict_class_c_w(X, weights):
    list = []
    for prob in predict_proba_c_w(X, weights): # Assigning range if <= 0.5 then assign t
        if prob <= 0.5:
            list.append(0)
        else:
            list.append(1)
    return list

```

In [87]:

```
def binary_loss_c_w(X, Y, weights, alpha=0.00001):
    l = -Y * np.log(predict_proba_c_w(X, weights)) - (1 - Y) * np.log(1 - predict_proba_c_w(X, weights))

    # L2 Regularisation
    l += alpha * weights.dot(weights)

    # Return the average Loss
    return np.mean(l)
```

In [88]:

```
def SGD_c_w(X_tr, Y_tr, X_dev, Y_dev, lr=0.1, alpha=0.00001, epochs=5, tolerance=0.0001, print_progress=True):
    np.random.seed(123)
    training_loss_history_c_w = []
    validation_loss_history_c_w = []

    weights_int_list_c_w = []
    for i in range(train_count_c_w.shape[1]):
        weights_int_list_c_w.append(0)
    weights_int_c_w = np.array(weights_int_list_c_w)
    weights_c_w = weights_int_c_w.astype(np.float)

    def zipper(X_tr, Y_tr):
        size = len(X_tr) if len(X_tr) < len(Y_tr) else len(Y_tr)
        retList = []
        for i in range(size):
            retList.append((X_tr[i], Y_tr[i]))
        return retList

    train_docs_c_w = zipper(X_tr, Y_tr)

    for epoch in range(epochs):
        np.random.shuffle(train_docs_c_w)

        for first, second in train_docs_c_w:
            weights_c_w = weights_c_w - lr * (first * (predict_proba_c_w(first, weights_c_w) - second))

        # Monitor training and validation loss
        loss_in_training_c_w = binary_loss_c_w(X_tr, Y_tr, weights_c_w, alpha)
        loss_in_dev_c_w = binary_loss_c_w(X_dev, Y_dev, weights_c_w, alpha)

        if epoch > 0 and validation_loss_history_c_w[-1] - loss_in_dev_c_w < tolerance:
            break
        else:
            training_loss_history_c_w.append(loss_in_training_c_w)
            validation_loss_history_c_w.append(loss_in_dev_c_w)

    if print_progress:
        print(f'Epoch: {epoch} | Training Loss: {loss_in_training_c_w} | Validation Loss: {loss_in_dev_c_w}')
    return weights_c_w, training_loss_history_c_w, validation_loss_history_c_w
```

```
w_count_c_w, training_loss_count_c_w, dev_loss_count_c_w = SGD_c_w(X_tr=train_count_c_w, Y_t
```

Epoch:- 0	Training loss:-	0.6867616762953799	Validation loss:-	0.688662994837779
Epoch:- 1	Training loss:-	0.6805971773460067	Validation loss:-	0.6843968844521979
Epoch:- 2	Training loss:-	0.6746274686149971	Validation loss:-	0.680144417580446
Epoch:- 3	Training loss:-	0.6688696422950303	Validation loss:-	0.6760317758642047
Epoch:- 4	Training loss:-	0.6633003554688951	Validation loss:-	0.6721490079854291
Epoch:- 5	Training loss:-	0.6579190436566391	Validation loss:-	0.6683357562193539
Epoch:- 6	Training loss:-	0.652708069921851	Validation loss:-	0.6647520226582013
Epoch:- 7	Training loss:-	0.6476875003339986	Validation loss:-	0.661629564605675
Epoch:- 8	Training loss:-	0.6427935428904774	Validation loss:-	0.6579495679975379
Epoch:- 9	Training loss:-	0.6380817207680346	Validation loss:-	0.6550194349984351
Epoch:- 10	Training loss:-	0.633490209212943	Validation loss:-	0.6516038096941142
Epoch:- 11	Training loss:-	0.6291039194471432	Validation loss:-	0.6484350345509996
Epoch:- 12	Training loss:-	0.6247632285196985	Validation loss:-	0.6457145349567553
Epoch:- 13	Training loss:-	0.6206123686762235	Validation loss:-	0.6431694524223667
Epoch:- 14	Training loss:-	0.616570545733435	Validation loss:-	0.6401596229146272
Epoch:- 15	Training loss:-	0.6126590560566499	Validation loss:-	0.6375004468692987
Epoch:- 16	Training loss:-	0.6088642492086557	Validation loss:-	0.6349291652259604
Epoch:- 17	Training loss:-	0.6051523149204795	Validation loss:-	0.6325877278185086
Epoch:- 18	Training loss:-	0.6015951747574431	Validation loss:-	0.6300874937335704
Epoch:- 19	Training loss:-	0.5980882283160485	Validation loss:-	0.6281110730962193
Epoch:- 20	Training loss:-	0.5946988075261816	Validation loss:-	0.6258516184296478
Epoch:- 21	Training loss:-	0.5914627284995453	Validation loss:-	0.6233718479773764
Epoch:- 22	Training loss:-	0.5882085807335026	Validation loss:-	0.6208516184296478

216552088358144

Epoch:- 23 Training loss:- 0.5850930112142171 Validation loss:- 0.6

195816096248331

Epoch:- 24 Training loss:- 0.5821583022108827 Validation loss:- 0.6

181383090807161

Epoch:- 25 Training loss:- 0.5791482098963159 Validation loss:- 0.6

155017473182208

Epoch:- 26 Training loss:- 0.5762623364447237 Validation loss:- 0.6

137055591790329

Epoch:- 27 Training loss:- 0.5734464949804937 Validation loss:- 0.6

120066411912459

Epoch:- 28 Training loss:- 0.5707129753882814 Validation loss:- 0.6

104406885078096

Epoch:- 29 Training loss:- 0.5681175705629719 Validation loss:- 0.6

091725443667702

Epoch:- 30 Training loss:- 0.5654810310599871 Validation loss:- 0.6

073639214180108

Epoch:- 31 Training loss:- 0.5629879815334035 Validation loss:- 0.6

052839427583832

Epoch:- 32 Training loss:- 0.5604983064944548 Validation loss:- 0.6

038306002210284

Epoch:- 33 Training loss:- 0.5580649621142231 Validation loss:- 0.6

02481923366284

Epoch:- 34 Training loss:- 0.5557124352510104 Validation loss:- 0.6

010701971209971

Epoch:- 35 Training loss:- 0.5534830319273188 Validation loss:- 0.6

002196303535654

Epoch:- 36 Training loss:- 0.5511746259711399 Validation loss:- 0.5

984922635513193

Epoch:- 37 Training loss:- 0.5489820218948812 Validation loss:- 0.5

971125866760734

Epoch:- 38 Training loss:- 0.5468850902229903 Validation loss:- 0.5

955819495972406

Epoch:- 39 Training loss:- 0.5447826686304875 Validation loss:- 0.5

943774214123583

Epoch:- 40 Training loss:- 0.5427061144831757 Validation loss:- 0.5

933492636292556

Epoch:- 41 Training loss:- 0.5407021400928198 Validation loss:- 0.5

922712453227861

Epoch:- 42 Training loss:- 0.5388685241466262 Validation loss:- 0.5

917774681444494

Epoch:- 43 Training loss:- 0.5368896478296042 Validation loss:- 0.5

897209666419778

Epoch:- 44 Training loss:- 0.5349776713666059 Validation loss:- 0.5

887529008488324

Epoch:- 45 Training loss:- 0.5331258667342416 Validation loss:- 0.5

880671738219826

Epoch:- 46 Training loss:- 0.531322620414796 Validation loss:- 0.58

70000258669776

Epoch:- 47 Training loss:- 0.5296009264937869 Validation loss:- 0.5

856664354604934

Epoch:- 48 Training loss:- 0.5278545684509905 Validation loss:- 0.5

847807360393954

Epoch:- 49 Training loss:- 0.5261447363643502 Validation loss:- 0.5

841580790826334

Epoch:- 50 Training loss:- 0.5244857371950931 Validation loss:- 0.5

832468943823458

Epoch:- 51 Training loss:- 0.5228839737238052 Validation loss:- 0.5

825608237935246

Epoch:- 52 Training loss:- 0.521277896509942 Validation loss:- 0.58

12450315946218



Epoch:- 53	Training loss:-	0.5197247895941459	Validation loss:-	0.5
80367357089618				
Epoch:- 54	Training loss:-	0.5181674734503027	Validation loss:-	0.5
796972510425905				
Epoch:- 55	Training loss:-	0.5166592372909627	Validation loss:-	0.5
789655192792105				
Epoch:- 56	Training loss:-	0.5152487292584949	Validation loss:-	0.5
787014311760628				
Epoch:- 57	Training loss:-	0.5137501020617173	Validation loss:-	0.5
777143269648036				
Epoch:- 58	Training loss:-	0.5123152713418281	Validation loss:-	0.5
768903779211736				
Epoch:- 59	Training loss:-	0.5109103009001279	Validation loss:-	0.5
76068126330121				
Epoch:- 60	Training loss:-	0.5095393746147446	Validation loss:-	0.5
754485778383727				
Epoch:- 61	Training loss:-	0.5081978798439779	Validation loss:-	0.5
745569078734861				
Epoch:- 62	Training loss:-	0.5068619296677492	Validation loss:-	0.5
740509213627891				
Epoch:- 63	Training loss:-	0.50555948026714	Validation loss:-	0.573
4506742333619				
Epoch:- 64	Training loss:-	0.5043157554149298	Validation loss:-	0.5
73154872632372				
Epoch:- 65	Training loss:-	0.5030267160153955	Validation loss:-	0.5
723492329141104				
Epoch:- 66	Training loss:-	0.5017850154809316	Validation loss:-	0.5
717171555470087				
Epoch:- 67	Training loss:-	0.5005655275200466	Validation loss:-	0.5
710728492805731				
Epoch:- 68	Training loss:-	0.49937266785066725	Validation loss:-	0.
5706135129869576				
Epoch:- 69	Training loss:-	0.4981951926942913	Validation loss:-	0.5
698506578387003				
Epoch:- 70	Training loss:-	0.49704698076252124	Validation loss:-	0.
5692451698035732				
Epoch:- 71	Training loss:-	0.49589635595008	Validation loss:-	0.568
8142794487164				
Epoch:- 72	Training loss:-	0.49477432247860914	Validation loss:-	0.
5683067830952089				
Epoch:- 73	Training loss:-	0.4936762189884289	Validation loss:-	0.5
681533642083794				
Epoch:- 74	Training loss:-	0.4925826898827807	Validation loss:-	0.5
676388011424371				
Epoch:- 75	Training loss:-	0.4915478038332081	Validation loss:-	0.5
674086923425322				
Epoch:- 76	Training loss:-	0.4904504826608226	Validation loss:-	0.5
666792008689355				
Epoch:- 77	Training loss:-	0.4894083185731051	Validation loss:-	0.5
661784494592169				
Epoch:- 78	Training loss:-	0.48838420457363463	Validation loss:-	0.
5656643851179889				
Epoch:- 79	Training loss:-	0.48737936493338685	Validation loss:-	0.
5651928738968522				

In [90]:

```
plt.plot(training_loss_count_c_w, label='Train_loss')
plt.plot(dev_loss_count_c_w, label='Valid_loss')

plt.xlabel('Epochs',fontweight='bold')
plt.ylabel('Loss',fontweight='bold')


plt.title('Binary_Count(Training)')
ax = plt.axes()
ax.set_facecolor("lightgray")

plt.legend()

plt.show()
```

<ipython-input-90-8bcc12968281>:8: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

```
ax = plt.axes()
```



In [91]:

```
array_test_label_c_w = np.array(test_label) # Changing list to array as accuracy score only
```

In [92]:

```
preds_te_count_c_w = predict_class_c_w(test_count_c_w, w_count_c_w)

print('Accuracy:', accuracy_score(array_test_label_c_w,preds_te_count_c_w))
print('Precision:', precision_score(array_test_label_c_w,preds_te_count_c_w))
print('Recall:', recall_score(array_test_label_c_w,preds_te_count_c_w))
print('F1-Score:', f1_score(array_test_label_c_w,preds_te_count_c_w))
```

```
Accuracy: 0.7192982456140351
Precision: 0.7164179104477612
Recall: 0.7236180904522613
F1-Score: 0.7200000000000001
```

In [93]:

```
w_tfidf_c_w, training_loss_tfidf_c_w, dev_loss_tfidf_c_w = SGD_c_w(X_tr=train_tfidf_c_w,
                                                                    Y_tr=np.array(train_label),
                                                                    X_dev=dev_tfidf_c_w,
                                                                    Y_dev=np.array(dev_label),
                                                                    lr=0.00322,
                                                                    alpha=0.0005,
                                                                    epochs=100)
```

```
Epoch:- 0    Training loss:- 0.6891974926424698    Validation loss:- 0.69
03373432508986
Epoch:- 1    Training loss:- 0.6853670496553458    Validation loss:- 0.68
76460321183888
Epoch:- 2    Training loss:- 0.6816495692504321    Validation loss:- 0.68
50171312064209
Epoch:- 3    Training loss:- 0.6780439691953213    Validation loss:- 0.68
24209757479942
```

<ipython-input-88-a7f79af50d5e>:11: DeprecationWarning: `np.float` is a deprecated alias for the builtin `float`. To silence this warning, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here.

Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations> (<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>)

```
weights_c_w = weights_int_c_w.astype(np.float)
```

```
Epoch:- 4    Training loss:- 0.6745456536671272    Validation loss:- 0.67
99691079814295
Epoch:- 5    Training loss:- 0.6711524108682186    Validation loss:- 0.67
75533859872013
Epoch:- 6    Training loss:- 0.667857922337171    Validation loss:- 0.675
2793165097853
Epoch:- 7    Training loss:- 0.6646632891414384    Validation loss:- 0.67
32508632111102
Epoch:- 8    Training loss:- 0.6615558623549975    Validation loss:- 0.67
09941230273697
Epoch:- 9    Training loss:- 0.6585467275733858    Validation loss:- 0.66
90670033026984
Epoch:- 10   Training loss:- 0.6556144743540093    Validation loss:- 0.6
669480704477625
Epoch:- 11   Training loss:- 0.6527868784207067    Validation loss:- 0.6
649095360123098
Epoch:- 12   Training loss:- 0.6500180807131299    Validation loss:- 0.6
631277366351428
Epoch:- 13   Training loss:- 0.6473404658174516    Validation loss:- 0.6
614914427873714
Epoch:- 14   Training loss:- 0.6447368374068324    Validation loss:- 0.6
59614712053002
Epoch:- 15   Training loss:- 0.6422114529019294    Validation loss:- 0.6
578943194553456
Epoch:- 16   Training loss:- 0.6397580318038434    Validation loss:- 0.6
562416007971652
Epoch:- 17   Training loss:- 0.6373656161643616    Validation loss:- 0.6
547108558669336
Epoch:- 18   Training loss:- 0.6350551546753361    Validation loss:- 0.6
5313107600437
Epoch:- 19   Training loss:- 0.6327881982145934    Validation loss:- 0.6
518195147935503
Epoch:- 20   Training loss:- 0.6305961951313848    Validation loss:- 0.6
```

504082519035101

Epoch:- 21 Training loss:- 0.6284788911851917 Validation loss:- 0.6

488664460731325

Epoch:- 22 Training loss:- 0.6263929774499469 Validation loss:- 0.6

476843055550392

Epoch:- 23 Training loss:- 0.6243789157811888 Validation loss:- 0.6

464236409064694

Epoch:- 24 Training loss:- 0.6224503005419547 Validation loss:- 0.6

454341081124544

Epoch:- 25 Training loss:- 0.6205156034235068 Validation loss:- 0.6

439030858602147

Epoch:- 26 Training loss:- 0.6186648264049013 Validation loss:- 0.6

427011563243432

Epoch:- 27 Training loss:- 0.6168563272364022 Validation loss:- 0.6

416093591020052

Epoch:- 28 Training loss:- 0.6150978103193234 Validation loss:- 0.6

40617987903677

Epoch:- 29 Training loss:- 0.6134072282900002 Validation loss:- 0.6

397628802456576

Epoch:- 30 Training loss:- 0.61173351849451 Validation loss:- 0.638

6843327495056

Epoch:- 31 Training loss:- 0.6101161411857254 Validation loss:- 0.6

374410005904216

Epoch:- 32 Training loss:- 0.6085435874709294 Validation loss:- 0.6

364841901741802

Epoch:- 33 Training loss:- 0.6070001918947184 Validation loss:- 0.6

356343026820224

Epoch:- 34 Training loss:- 0.6055048167329065 Validation loss:- 0.6

347605338144376

Epoch:- 35 Training loss:- 0.6040616191587256 Validation loss:- 0.6

341563531303621

Epoch:- 36 Training loss:- 0.6026316388706937 Validation loss:- 0.6

332097299765354

Epoch:- 37 Training loss:- 0.6012493849794215 Validation loss:- 0.6

323750884907592

Epoch:- 38 Training loss:- 0.5999145657063153 Validation loss:- 0.6

314496183118206

Epoch:- 39 Training loss:- 0.5986067765610381 Validation loss:- 0.6

306907762318517

Epoch:- 40 Training loss:- 0.597318994272567 Validation loss:- 0.63

00572928950062

Epoch:- 41 Training loss:- 0.5960723145216986 Validation loss:- 0.6

294253455222294

Epoch:- 42 Training loss:- 0.5948915617894784 Validation loss:- 0.6

290337520179067

Epoch:- 43 Training loss:- 0.5936873521815146 Validation loss:- 0.6

27973239541002

Epoch:- 44 Training loss:- 0.5925313512370677 Validation loss:- 0.6

273692610200647

Epoch:- 45 Training loss:- 0.5914029641017373 Validation loss:- 0.6

269226767494822

Epoch:- 46 Training loss:- 0.5903065141183284 Validation loss:- 0.6

263243255049634

Epoch:- 47 Training loss:- 0.5892495104195308 Validation loss:- 0.6

255779417807289

Epoch:- 48 Training loss:- 0.5882036370170088 Validation loss:- 0.6

25054354596652

Epoch:- 49 Training loss:- 0.5871784602377996 Validation loss:- 0.6

24682602242702

Epoch:- 50 Training loss:- 0.5861865865885991 Validation loss:- 0.6

241775953024264

Epoch:- 51	Training loss:- 0.5852271342319111	Validation loss:- 0.6
23795311802667		
Epoch:- 52	Training loss:- 0.5842773418515842	Validation loss:- 0.6
231023077972939		
Epoch:- 53	Training loss:- 0.583360780710365	Validation loss:- 0.62
26047462931033		
Epoch:- 54	Training loss:- 0.5824549524376215	Validation loss:- 0.6
222337196654332		
Epoch:- 55	Training loss:- 0.5815767382920334	Validation loss:- 0.6
218292053620608		
Epoch:- 56	Training loss:- 0.5807387313868261	Validation loss:- 0.6
216619565996286		
Epoch:- 57	Training loss:- 0.5798926487958017	Validation loss:- 0.6
211919585072774		
Epoch:- 58	Training loss:- 0.5790719449309535	Validation loss:- 0.6
207619047925461		
Epoch:- 59	Training loss:- 0.5782711918190223	Validation loss:- 0.6
203267340179767		
Epoch:- 60	Training loss:- 0.5774936585901443	Validation loss:- 0.6
199907509757269		
Epoch:- 61	Training loss:- 0.5767366889201465	Validation loss:- 0.6
195256211698039		
Epoch:- 62	Training loss:- 0.575990672558028	Validation loss:- 0.61
92869784736128		
Epoch:- 63	Training loss:- 0.5752659068957235	Validation loss:- 0.6
18977663544562		
Epoch:- 64	Training loss:- 0.5745730196538625	Validation loss:- 0.6
188618123923215		
Epoch:- 65	Training loss:- 0.5738707373205054	Validation loss:- 0.6
184784401732966		
Epoch:- 66	Training loss:- 0.5731919719340509	Validation loss:- 0.6
181552204786027		
Epoch:- 67	Training loss:- 0.5725317577554059	Validation loss:- 0.6
178756421612964		
Epoch:- 68	Training loss:- 0.5718878117764333	Validation loss:- 0.6
176387959442443		
Epoch:- 69	Training loss:- 0.5712561688729434	Validation loss:- 0.6
172564103770853		
Epoch:- 70	Training loss:- 0.570646006965033	Validation loss:- 0.61
6951171985971		
Epoch:- 71	Training loss:- 0.5700411505974223	Validation loss:- 0.6
167477685043423		
Epoch:- 72	Training loss:- 0.5694526498515728	Validation loss:- 0.6
165222922624678		

In [94]:

```
plt.plot(training_loss_tfidf_c_w, label='Train_loss')
plt.plot(dev_loss_tfidf_c_w, label='Valid_loss')

plt.xlabel('Epochs', fontweight='bold')
plt.ylabel('Loss', fontweight='bold')


plt.title('Binary - TFIDF(Training)')
ax = plt.axes()
ax.set_facecolor("lightgray")

plt.legend()

plt.show()
```

<ipython-input-94-2895b5c1aa72>:8: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

```
ax = plt.axes()
```



In [95]:

```
preds_te_tfidf_c_w = predict_class_c_w(test_tfidf_c_w, w_tfidf_c_w)

print('Accuracy:', accuracy_score(np.array(test_label),preds_te_tfidf_c_w))
print('Precision:', precision_score(np.array(test_label),preds_te_tfidf_c_w))
print('Recall:', recall_score(np.array(test_label),preds_te_tfidf_c_w))
print('F1-Score:', f1_score(np.array(test_label),preds_te_tfidf_c_w))
```

Accuracy: 0.7218045112781954

Precision: 0.72

Recall: 0.7236180904522613

F1-Score: 0.7218045112781956

## Full Results

Add here your results:

LR	Precision	Recall	F1-Score
BOW-count	0.8454106280193237	0.8793969849246231	0.8620689655172413
BOW-tfidf	0.8514851485148515	0.864321608040201	0.85785536159601
BOCN-count	0.7046632124352331	0.6834170854271356	0.6938775510204082
BOCN-tfidf	0.7114427860696517	0.7185929648241206	0.7150000000000001
BOW+BOCN-count	0.7164179104477612	0.7236180904522613	0.7200000000000001
BOW+BOCN-tfidf	0.72	0.7236180904522613	0.7218045112781956

Please discuss why your best performing model is better than the rest.

My model is having a great accuracy score, precision score, Recall score and F1-score in all three sections which implies that it is predicting better results and giving output at its best. Thats why I think my model is performing better than the rest.